

Assignment1

June 28, 2023

1 Assignment 1

1.1 Description

The census-income dataset contains census information for 48,842 people. It has 14 attributes for each person (age, workclass, fnlwgt, education, education-num, marital-status, occupation, relationship, race, sex, capital-gain, capital-loss, hours-per-week, and native-country) and a Boolean attribute class classifying the input of the person as belonging to one of two categories >50K, <=50K. The prediction problem here is to classify whether a person's salary is >50K or <= 50K given the attribute values. ## Properties of Data - Number of Instances - 48842 instances, mix of continuous and discrete (train=32561, test=16281) - 45222 if instances with unknown values are removed (train=30162, test=15060) - Number of Attributes: 6 continuous, 8 nominal attributes - Attribute Information: 1) age: continuous 2) workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked 3) fnlwgt: continuous 4) education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool 5) education-num: continuous 6) marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse 7) occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces 8) relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried 9) race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black 10) sex: Female, Male 11) capital-gain: continuous 12) capital-loss: continuous 13) hours-per-week: continuous 14) native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands 15) class: >50K, <=50K

To get started with the model, run environment.yml file to initialize the conda environment with relevant libraries

1.2 Processing Training Data

```
[4]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
# import chefboost.Chefboost as chef
# import chefboost as chf
```

```

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn import tree
from sklearn.tree import plot_tree
import json
import pickle
import graphviz
%matplotlib inline

```

```

[5]: train_df = pd.read_csv("data1/train.csv", na_values="?")
test_df = pd.read_csv("data1/test.csv", na_values="?")
train_df[25:30]

```

```

[5]:
   age  workclass  fnlwgt  education  education-num  marital-status \
25   56   Local-gov  216851   Bachelors             13  Married-civ-spouse
26   19    Private  168294    HS-grad              9    Never-married
27   54         NaN  180211  Some-college          10  Married-civ-spouse
28   39    Private  367260    HS-grad              9        Divorced
29   49    Private  193366    HS-grad              9  Married-civ-spouse

   occupation  relationship      race  sex  capital-gain \
25   Tech-support      Husband    White  Male           0
26   Craft-repair  Own-child    White  Male           0
27         NaN      Husband  Asian-Pac-Islander  Male           0
28  Exec-managerial  Not-in-family    White  Male           0
29   Craft-repair      Husband    White  Male           0

   capital-loss  hours-per-week  native-country  class
25           0             40  United-States  >50K
26           0             40  United-States  <=50K
27           0             60         South  >50K
28           0             80  United-States  <=50K
29           0             40  United-States  <=50K

```

```

[4]: train_df['workclass'].value_counts()

```

```

[4]: Private           22696
Self-emp-not-inc     2541
Local-gov            2093
State-gov            1298
Self-emp-inc         1116
Federal-gov          960
Without-pay          14
Never-worked         7
Name: workclass, dtype: int64

```

```
[6]: mode_values = train_df.median(numeric_only=True)
train_df.fillna(mode_values, inplace=True)
mode_values
```

```
[6]: age                37.0
fnlwgt            178356.0
education-num      10.0
capital-gain        0.0
capital-loss        0.0
hours-per-week     40.0
dtype: float64
```

```
[7]: mode_values = []
for column in train_df.columns:
    if train_df[column].dtype == 'object':
        mode_value = train_df[column].mode()[0]
        mode_values.append(mode_value)
        train_df[column].fillna(mode_value, inplace=True)
mode_values
```

```
[7]: ['Private',
      'HS-grad',
      'Married-civ-spouse',
      'Prof-specialty',
      'Husband',
      'White',
      'Male',
      'United-States',
      '<=50K']
```

```
[8]: train_df.rename(columns={'class': 'Decision'}, inplace=True)
test_df.rename(columns={'class': 'Decision'}, inplace=True)
train_df[25:30]
```

```
[8]:
```

	age	workclass	fnlwgt	education	education-num	marital-status	\
25	56	Local-gov	216851	Bachelors	13	Married-civ-spouse	
26	19	Private	168294	HS-grad	9	Never-married	
27	54	Private	180211	Some-college	10	Married-civ-spouse	
28	39	Private	367260	HS-grad	9	Divorced	
29	49	Private	193366	HS-grad	9	Married-civ-spouse	

	occupation	relationship	race	sex	capital-gain	\
25	Tech-support	Husband	White	Male	0	
26	Craft-repair	Own-child	White	Male	0	
27	Prof-specialty	Husband	Asian-Pac-Islander	Male	0	
28	Exec-managerial	Not-in-family	White	Male	0	
29	Craft-repair	Husband	White	Male	0	

	capital-loss	hours-per-week	native-country	Decision
25	0	40	United-States	>50K
26	0	40	United-States	<=50K
27	0	60	South	>50K
28	0	80	United-States	<=50K
29	0	40	United-States	<=50K

We have finally converted our training data to be used for training, by replacing NaN values with most repeated values for the respective columns, medians for numerical data, and converted the final class into a column called “Decision”.

One quirk in chefboost is the approach to the *target variable*— it must be stored in the same dataframe as the features, it must be called Decision and must be the very last column of the dataframe. Quite weird, but there is probably some good reason for that.

1.3 Handling continuous values

For this algorithm, we’re using Chefboost, an algorithm which allows us to implement C4.5 algorithm for Decision Tree, and allows us to calculate gain ratio as well. We will use this data to find points of discretization and make continuous values into discrete data, which can be fed to our machine.

```
[31]: #!/pip install chefboost
from chefboost.training import Training
config = {'algorithm': 'C4.5'}
```

```
[38]: def gainratiocal(threshold:float, column:object) -> float:
    idx = train_df[train_df[f"{column}"] <= threshold].index
    tmp_df = train_df.copy()
    tmp_df[f"{column}"] = f">{threshold}"
    tmp_df.loc[idx, f"{column}"] = f"<={threshold}"
    grat = Training.findGains(tmp_df, config)['gains'][f"{column}"]
    return grat
```

```
[39]: for i in range(1,100,10):
    a = gainratiocal(float(i), 'hours-per-week')
    print(i,":",a)
```

```
1 : 0.01753148492873723
11 : 0.017751847016788638
21 : 0.018828987795244542
31 : 0.02064893040902423
41 : 0.023904627769250456
51 : 0.017662187191664586
61 : 0.008553549138154704
71 : 0.007151079969351616
81 : 0.00477984617092674
91 : 0.0023187326063295326
```

The above data clearly calculates the gain ratios for different values, from which we can clearly see that data around 31 and 41 may give desired decision boundary

1.3.1 Other calculations

Here, we've done similar calculation for other remaining continuous attributes

```
[43]: for i in range (35, 51,2):  
        a = gainratiocal(float(i), 'hours-per-week')  
        print(i,":",a)
```

```
35 : 0.021594639037470862  
37 : 0.02197617790402921  
39 : 0.022415767382654897  
41 : 0.023904627769250456  
43 : 0.02476910237843594  
45 : 0.024482106814724198  
47 : 0.024970789695430925  
49 : 0.026858965904353535
```

We can observe that 49 would be the value with the highest gain ratio, giving us the decision boundary for this continuous column, we can do similar calculation for the rest, which is done so by chefboost model internally.

```
[47]: edunum = sorted(train_df['education-num'].unique())
```

```
[49]: for i in edunum:  
        a = gainratiocal(float(i), 'education-num')  
        print(i,":",a)
```

```
1 : 0.03192754615584799  
2 : 0.03196053650702954  
3 : 0.0321292294421387  
4 : 0.032515666540499984  
5 : 0.03304882274478729  
6 : 0.03394432885606893  
7 : 0.03527948364873703  
8 : 0.03610098814658912  
9 : 0.03933489231796834  
10 : 0.05295739290033519  
11 : 0.0604007911744962  
12 : 0.06955873321174283  
13 : 0.08772951926659586  
14 : 0.1099725423325889  
15 : 0.10501869293424891  
16 : 0.0
```

For education-num, the value 14 establishes a good decision boundary

```
[56]: age = sorted(train_df['age'].unique())
      for i in range(25,30):
          a = gainratiocal(float(i), 'age')
          print(i,":",a)
```

```
25 : 0.018878608763892753
26 : 0.018949020158395617
27 : 0.019008258539748947
28 : 0.018828616217681778
29 : 0.018589001236071493
```

For age, 27 is the decision boundary

```
[77]: fnlwgt = sorted(train_df["fnlwgt"].unique())
      for i in range(12200, 22200, 1000):
          a = gainratiocal(float(i), 'fnlwgt')
          print(i,":",a)
```

```
12200 : 0.04011843427056207
13200 : 0.04011843427056207
14200 : 0.04011860831741389
15200 : 0.04011287222370235
16200 : 0.04011287222370235
17200 : 0.04011287222370235
18200 : 0.04011287222370235
19200 : 0.04011209119723716
20200 : 0.04009173689623385
21200 : 0.04006503030840085
```

For fnlwgt, the boundary can be taken around 14200, although we may take more than one boundary for this case

1.4 Training the Chefboost Model

For the C4.5 decision tree, we will use Chefboost library, which provides the model we require. We choose this light-weight library because: - support of categorical features, meaning we do not need to pre-process them using, for example, one-hot encoding. - the decision trees trained using chefboost are stored as if-else statements in a dedicated Python file. This way, we can easily see what decisions the tree makes to arrive at a given prediction. - we can choose one of the multiple algorithms to train the decision trees.

Chefboost implements different algorithms like ID3, C4.5, CART(this is implemented in scikit-learn). It also provides the models for random forest, which is why it is considered a better library for our particular assignment.

```
[8]: config = {'algorithm':'C4.5'}
      model = chef.fit(train_df, config)
```

```
[INFO]: 4 CPU cores will be allocated in parallel running
C4.5 tree is going to be built...
```

```
-----
```

finished in 2088.6931281089783 seconds


Evaluate train set

Accuracy: 88.01019624704401 % on 32561 instances
Labels: ['<=50K' '>50K']
Confusion matrix: [[23115, 2299], [1605, 5542]]
Precision: 90.9538 %, Recall: 93.5073 %, F1: 92.2129 %

```
[9]: rules = "outputs/rules/rules.py"  
fi = chef.feature_importance(rules).set_index("feature")  
fi.plot(kind="barh", title="Feature Importance")
```

Decision rule: outputs/rules/rules.py

```
[9]: <Axes: title={'center': 'Feature Importance'}, ylabel='feature'>
```



```
[4]: chfmdl = chef.load_model("model.pkl")
```

```
[9]: chef.evaluate(chfmdl, test_df)
```

Evaluate test set

Accuracy: 82.07112585222038 % on 16281 instances

```
Labels:  ['<=50K' '>50K']
Confusion matrix:  [[11136, 1620], [1299, 2226]]
Precision:  87.3001 %, Recall:  89.5537 %, F1:  88.4125 %
```

Here, we have found the entire tree, and the tree's dependence on different attributes. We will now begin plotting error for training and testing data, and finally obtain a tree of better accuracy. Now, due to the shortcomings of the `chebboost`, we really can't apply pruning, and hence will use CART algorithm, due to the lack of availability of C4.5 algorithm.

1.5 Training Scikit Learn model

To evaluate the final data and test different errors, we'll train the scikit decision tree model which works on CART algorithm, a much advanced algorithm

1.5.1 Preprocessing data

```
[58]: train_df.columns
```

```
[58]: Index(['age', 'fnlwgt', 'education-num', 'capital-gain', 'capital-loss',
          'hours-per-week', 'workclass_n', 'education_n', 'marital-status_n',
          'occupation_n', 'relationship_n', 'race_n', 'sex_n', 'native-country_n',
          'class_n'],
          dtype='object')
```

```
[57]: train_df.select_dtypes(include=['object']).iloc[0]
```

```
[57]: Series([], Name: 0, dtype: float64)
```

These are the columns which need to be encoded, for which we will use label encoder, which allows us to perform one hot encoding.

```
[41]: le_workc = LabelEncoder()
      le_educ = LabelEncoder()
      le_mari = LabelEncoder()
      le_occup = LabelEncoder()
      le_relat = LabelEncoder()
      le_race = LabelEncoder()
      le_sex = LabelEncoder()
      le_native = LabelEncoder()
      le_class = LabelEncoder()
```

```
[42]: train_df["workclass_n"] = le_workc.fit_transform(train_df["workclass"])
      train_df["education_n"] = le_educ.fit_transform(train_df["education"])
      train_df["marital-status_n"] = le_mari.fit_transform(train_df["marital-status"])
      train_df["occupation_n"] = le_occup.fit_transform(train_df["occupation"])
      train_df["relationship_n"] = le_relat.fit_transform(train_df["relationship"])
      train_df["race_n"] = le_race.fit_transform(train_df["race"])
      train_df["sex_n"] = le_sex.fit_transform(train_df["sex"])
```



```

train_df["native-country_n"] = le_native.
↳fit_transform(train_df["native-country"])
train_df["class_n"] = le_class.fit_transform(train_df["class"])
train_df.iloc[0]

```

```

[42]: age                39
      workclass          State-gov
      fnlwgt             77516
      education          Bachelors
      education-num       13
      marital-status      Never-married
      occupation          Adm-clerical
      relationship        Not-in-family
      race                White
      sex                 Male
      capital-gain         2174
      capital-loss         0
      hours-per-week       40
      native-country      United-States
      class                <=50K
      workclass_n          6
      education_n          9
      marital-status_n      4
      occupation_n         0
      relationship_n        1
      race_n               4
      sex_n                1
      native-country_n      38
      class_n              0
      Name: 0, dtype: object

```

Now that we have successfully encoded the classes, we can drop the unwanted columns

```

[60]: columns=['class', 'native-country', 'sex', 'race', 'relationship', 'occupation', 'marital-status', 'e
X_train = train_df.drop(columns=columns,axis='columns')
X_train.head()

```

```

[60]:      age  fnlwgt  education-num  capital-gain  capital-loss  hours-per-week  \
39203   39  265685           10           0           0           65
16702   49  281647           14           0           0           45
43825   19  163885            9           0           0           40
48735   64   47298           16           0           0           45
34480   46  146786            6           0           0           50

      workclass_n  education_n  marital-status_n  occupation_n  \
39203           8.0          15                0           14.0
16702           3.0          12                2            3.0
43825           3.0          11                4            6.0

```

48735	1.0	10	2	10.0
34480	3.0	0	0	13.0

	relationship_n	race_n	sex_n	native-country_n
39203	1	4	1	32.0
16702	0	4	1	38.0
43825	3	4	0	38.0
48735	0	4	1	38.0
34480	1	4	1	38.0

We have successfully encoded training data into X, and can similarly do that for Y, and train our model. Also, we'll retrieve the mappings of each individual, and store them offline, to convert our test data when needed.

```
[61]: Y_train = train_df.class_n
      Y_train.head()
```

```
[61]: 0    0
      1    0
      2    0
      3    0
      4    0
      Name: class_n, dtype: int64
```

```
[ ]: label_mapping = {}
      label_mapping['workclass'] = {label: encoded_label for label, encoded_label in
      ↪ zip(train_df['workclass'], train_df['workclass_n'])}
      label_mapping
```

```
[51]: columns.remove('class_n')
      for column in columns:
          label_mapping[column] = {label: encoded_label for label, encoded_label in
          ↪ zip(train_df[column], train_df[column+'_n'])}
      label_mapping['sex']
```

```
[51]: {'Male': 1, 'Female': 0}
```

```
[54]: with open('label_mapping.json', 'w') as output_file:
      json.dump(label_mapping, output_file)
```

```
[7]: def encode_data_with_mapping(data_df, label_mapping_file='label_mapping.json'):
      # Load the label encoding mapping from the JSON file
      with open(label_mapping_file, 'r') as mapping_file:
          label_mapping = json.load(mapping_file)

      # Create a new DataFrame to store the encoded data
      encoded_data_df = pd.DataFrame()
```

```

for column in data_df.columns:
    if column not in label_mapping:
        encoded_data_df[column] = data_df[column]

# Iterate over the columns of the input DataFrame
for column in data_df.columns:
    if column in label_mapping:
        # Create a label encoder and set the classes using the label mapping
        label_encoder = LabelEncoder()
        label_encoder.classes_ = np.append(list(label_mapping[column].
↪values()), 'Unknown')

        # Apply the label encoding to the selected column
        encoded_column = data_df[column].astype(str).fillna('Unknown').
↪map(label_mapping[column]).fillna(len(label_mapping[column]))
        encoded_data_df[column+'_n'] = encoded_column

return encoded_data_df

```

We have now defined a function which will automatically do all of this for us

```

[11]: train_df = pd.read_csv("data1/train.csv")
train_df = encode_data_with_mapping(train_df)
train_df.head()

```

```

[11]:
  age  fnlwgt  education-num  capital-gain  capital-loss  hours-per-week  \
0   39   77516             13           2174             0             40
1   50   83311             13              0             0             13
2   38  215646              9              0             0             40
3   53  234721              7              0             0             40
4   28  338409             13              0             0             40

  workclass_n  education_n  marital-status_n  occupation_n  relationship_n  \
0           6.0           9                4           0.0             1
1           5.0           9                2           3.0             0
2           3.0          11                0           5.0             1
3           3.0           1                2           5.0             0
4           3.0           9                2           9.0             5

  race_n  sex_n  native-country_n  class_n
0       4     1             38.0         0
1       4     1             38.0         0
2       4     1             38.0         0
3       2     1             38.0         0
4       2     0              4.0         0

```

```
[8]: train_df = pd.read_csv("data1/train.csv", na_values="?")
train_df = encode_data_with_mapping(train_df)
X_train = train_df.drop(columns='class_n', axis='columns')
Y_train = train_df.class_n
```

```
[9]: test_df = pd.read_csv("data1/test.csv", na_values="?")
test_df = encode_data_with_mapping(test_df)
test_df.head()
```

```
[9]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	\
0	25	226802	7	0	0	40	
1	38	89814	9	0	0	50	
2	28	336951	12	0	0	40	
3	44	160323	10	7688	0	40	
4	18	103497	10	0	0	30	

	workclass_n	education_n	marital-status_n	occupation_n	relationship_n	\
0	3.0	1	4	6.0	3	
1	3.0	11	2	4.0	0	
2	1.0	7	2	10.0	0	
3	3.0	15	2	6.0	0	
4	8.0	15	4	14.0	3	

	race_n	sex_n	native-country_n	class_n
0	2	1	38.0	0
1	4	1	38.0	0
2	4	1	38.0	1
3	2	1	38.0	1
4	4	0	38.0	0

```
[10]: X_test = test_df.drop(columns='class_n', axis='columns')
Y_test = test_df.class_n
X_valid, X_test, Y_valid, Y_test = train_test_split(X_test, Y_test, test_size=0.
↪5, random_state=10)
X_valid.shape, X_test.shape
```

```
[10]: ((8140, 14), (8141, 14))
```

We have successfully split the test set into testing and validating dataset

1.5.2 Training and scoring

```
[67]: cartree = tree.DecisionTreeClassifier(criterion='entropy')
cartree.fit(X_train, Y_train)
```

```
[67]: DecisionTreeClassifier(criterion='entropy')
```

```
[68]: cartree.tree_.node_count
```

```
[68]: 9431
```

```
[69]: cartree.get_depth()
```

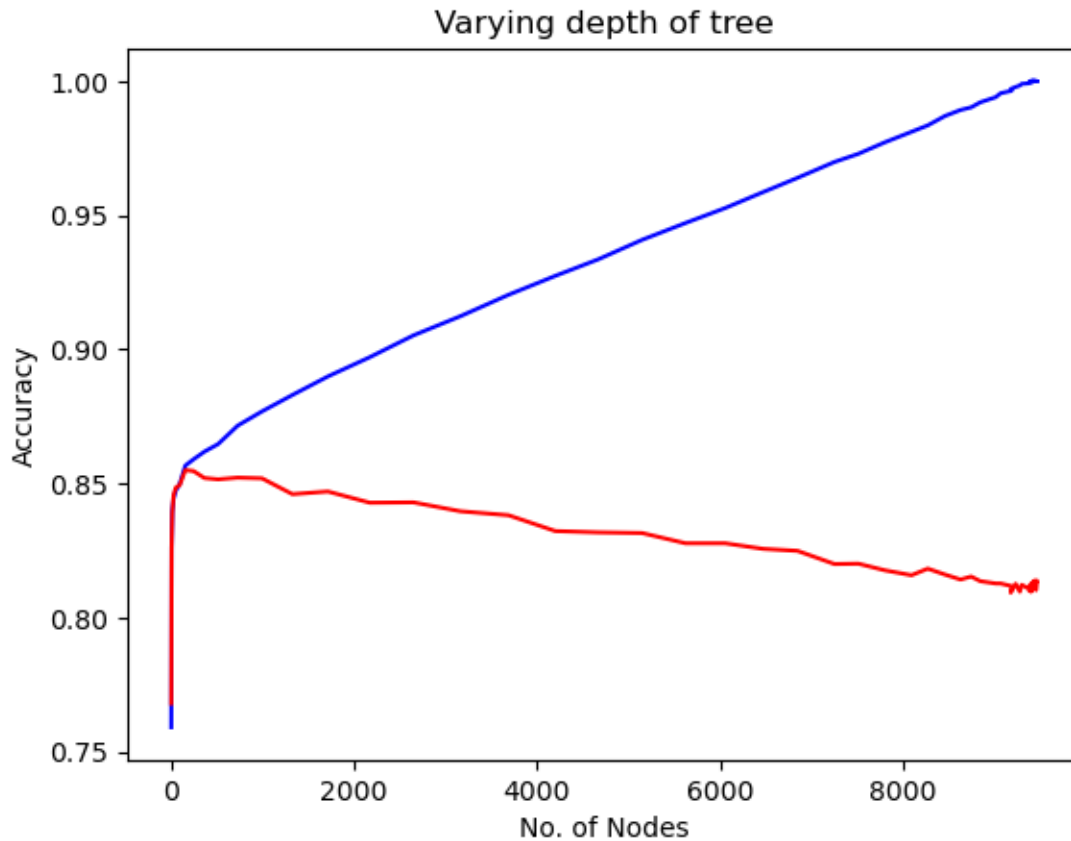
```
[69]: 49
```

Now, we can run a loop, restricting the depth to collect no of nodes and accuracy for both test and valid data, finally create a data for graphing the discrepancy caused by overfitting, as well as later pruning our resultant model.

```
[73]: trainscore = []
      testscore = []
      nofnodes = []
      depthcou = []
      for i in range(1, 55):
          clf = tree.DecisionTreeClassifier(criterion = 'entropy', max_depth=i)
          clf.fit(X_train, Y_train)
          trainscore.append(clf.score(X_train, Y_train))
          testscore.append(clf.score(X_valid, Y_valid))
          nofnodes.append(clf.tree_.node_count)
          depthcou.append(i)
```

```
[74]: plt.plot(nofnodes, trainscore, 'b')
      plt.plot(nofnodes, testscore, 'r')
      plt.xlabel("No. of Nodes")
      plt.ylabel("Accuracy")
      plt.title("Varying depth of tree")
```

```
[74]: Text(0.5, 1.0, 'Varying depth of tree')
```



```
[75]: maxacc = max(testscore)
      inx = testscore.index(maxacc)
      maxnodes = nofnodes[inx]
      maxnodes, depthcou[inx]
```


[75]: (157, 7)

Doing the same, but this time restricting no of nodes

```
[19]: trainscore1 = []
      testscore1 = []
      nofnodes1 = []
      noofleaf = []
      for i in range(2, 9503, 100):
          clf = tree.DecisionTreeClassifier(criterion = 'entropy', max_leaf_nodes=i)
          clf.fit(X_train, Y_train)
          trainscore1.append(clf.score(X_train, Y_train))
          testscore1.append(clf.score(X_valid, Y_valid))
          nofnodes1.append(clf.tree_.node_count)
          noofleaf.append(i)
```

```
[21]: plt.plot(nofnodes1, trainscore1, 'b')
plt.plot(nofnodes1, testscore1, 'r')
plt.xlabel("No. of Nodes")
plt.ylabel("Accuracy")
plt.title("Varying no. of Leaf nodes")
```

```
[21]: Text(0.5, 1.0, 'Varying no. of Leaf nodes')
```



```
[22]: maxacc = max(testscore1)
inx = testscore1.index(maxacc)
maxnodes = nofnodes1[inx]
maxleaf = noofleaf[inx]
maxnodes, maxleaf
```

```
[22]: (203, 102)
```

Taking average of two maxnodes value gives max nodes of 225

```
[14]: cartree = tree.DecisionTreeClassifier(criterion='entropy', max_depth=8)
cartree.fit(X_train, Y_train)
```

```
cartree.score(X_valid, Y_valid)
```

```
[14]: 0.8547911547911548
```


```
[113]: cartree.score(X_test, Y_test)
```

```
[113]: 0.8577570323056135
```

```
[15]: with open("DecisionTreeFinal.pkl", 'wb') as f:  
      pickle.dump(cartree, f)
```

We've got an accuracy of 85.8% for test data, which closely matches that of validation data

```
[22]: fi = cartree.feature_importances_  
fig, ax = plt.subplots()  
ax.barh(X_train.columns, fi)  
ax.invert_yaxis()
```



```
[23]: fig.savefig('FeatureDF.png')
```

```
[21]: with open("DecisionTreeFinal.pkl", 'rb') as f:  
      cartree = pickle.load(f)  
dot_data = tree.export_graphviz(cartree, out_file=None, feature_names=X_train.  
    ↪ columns, class_names=['<50K', '>=50K'], filled=True)
```



```

# Create a graph from the Graphviz data
graph = graphviz.Source(dot_data)

# Render the graph
graph.render("decision_tree")

# Display the graph
graph

```

[21]:



```

[17]: graph.format = 'png'
graph.render('DecisionFinal')

```

[17]: 'DecisionFinal.png.png'

1.6 Choosing random data and training

1.6.1 Preprocessing data

```

[28]: data = pd.read_csv("data1/combined.csv", na_values="?")
mode_values = data.median(numeric_only=True)
data.fillna(mode_values, inplace=True)
mode_values = []
for column in train_df.columns:
    if train_df[column].dtype == 'object':
        mode_value = train_df[column].mode()[0]
        mode_values.append(mode_value)
        train_df[column].fillna(mode_value, inplace=True)
data.head()

```

```

[28]:  age      workclass  fnlwgt  education  education-num  \
0    39      State-gov   77516    Bachelors             13
1    50  Self-emp-not-inc  83311    Bachelors             13
2    38      Private  215646    HS-grad              9
3    53      Private  234721      11th              7
4    28      Private  338409    Bachelors             13

   marital-status  occupation  relationship  race  sex  \
0  Never-married  Adm-clerical  Not-in-family  White  Male
1  Married-civ-spouse  Exec-managerial      Husband  White  Male
2      Divorced  Handlers-cleaners  Not-in-family  White  Male
3  Married-civ-spouse  Handlers-cleaners      Husband  Black  Male

```

	Married-civ-spouse	Prof-specialty	Wife	Black	Female
	capital-gain	capital-loss	hours-per-week	native-country	class
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

```
[29]: data = encode_data_with_mapping(data)
X = data.drop(columns='class_n', axis='columns')
Y = data.class_n
data.head()
```

```
[29]:
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week	\
0	39	77516	13	2174	0	40	
1	50	83311	13	0	0	13	
2	38	215646	9	0	0	40	
3	53	234721	7	0	0	40	
4	28	338409	13	0	0	40	

	workclass_n	education_n	marital-status_n	occupation_n	relationship_n	\
0	6.0	9	4	0.0	1	
1	5.0	9	2	3.0	0	
2	3.0	11	0	5.0	1	
3	3.0	1	2	5.0	0	
4	3.0	9	2	9.0	5	

	race_n	sex_n	native-country_n	class_n
0	4	1	38.0	0
1	4	1	38.0	0
2	4	1	38.0	0
3	2	1	38.0	0
4	2	0	4.0	0

```
[30]: X_train,X_test,Y_train, Y_test = train_test_split(X,Y,test_size=0.
↳33,random_state=10)
X_train.shape,X_test.shape
```

```
[30]: ((32724, 14), (16118, 14))
```

```
[31]: X_valid, X_test, Y_valid, Y_test = train_test_split(X_test, Y_test, test_size=0.
↳5, random_state=10)
X_valid.shape, X_test.shape
```

```
[31]: ((8059, 14), (8059, 14))
```

1.6.2 Training and scoring

```
[42]: cartree = tree.DecisionTreeClassifier(criterion='entropy')
      cartree.fit(X_train, Y_train)
```

```
[42]: DecisionTreeClassifier(criterion='entropy')
```

```
[43]: cartree.tree_.node_count
```

```
[43]: 9365
```

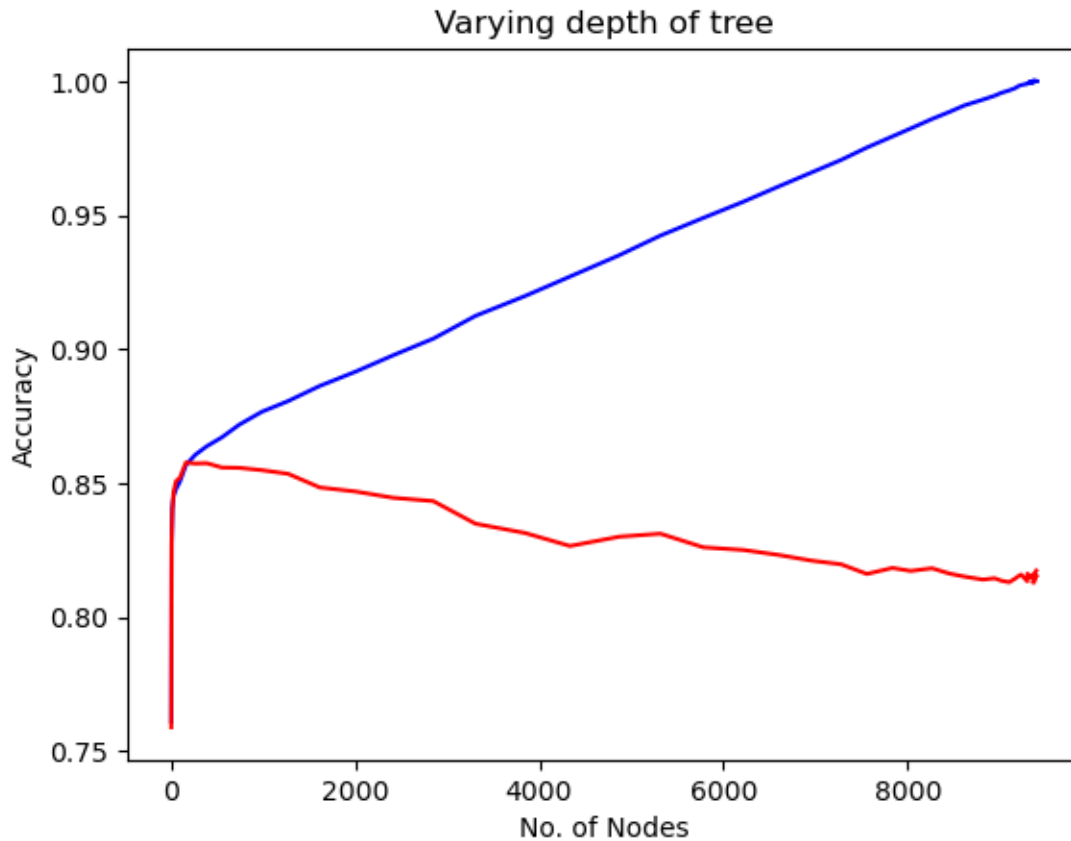
```
[44]: cartree.get_depth()
```

```
[44]: 48
```

```
[45]: trainscore = []
      testscore = []
      nofnodes = []
      depthcon = []
      for i in range(1, 55):
          clf = tree.DecisionTreeClassifier(criterion = 'entropy', max_depth=i)
          clf.fit(X_train, Y_train)
          trainscore.append(clf.score(X_train, Y_train))
          testscore.append(clf.score(X_valid, Y_valid))
          nofnodes.append(clf.tree_.node_count)
          depthcon.append(i)
```

```
[46]: plt.plot(nofnodes, trainscore, 'b')
      plt.plot(nofnodes, testscore, 'r')
      plt.xlabel("No. of Nodes")
      plt.ylabel("Accuracy")
      plt.title("Varying depth of tree")
```

```
[46]: Text(0.5, 1.0, 'Varying depth of tree')
```



```
[47]: maxacc = max(testscore)
      inx = testscore.index(maxacc)
      maxnodes = nofnodes[inx]
      maxnodes, depthcon[inx]
```

```
[47]: (163, 7)
```

```
[106]: cartree = tree.DecisionTreeClassifier(criterion='entropy', max_depth=8)
      cartree.fit(X_train, Y_train)
      cartree.score(X_valid, Y_valid)
```


```
[106]: 0.8554054054054054
```

```
[107]: cartree.score(X_test, Y_test)
```

```
[107]: 0.8577570323056135
```

```
[108]: with open("DecisionTreeRandom.pkl", 'wb') as f:
      pickle.dump(cartree, f)
```

```
[19]: fi1 = cartree.feature_importances_
fig, ax = plt.subplots()
ax.barh(X_train.columns, fi1)
ax.invert_yaxis()
```



```
[20]: fig.savefig('FeatureDR.png')
```

```
[18]: with open("DecisionTreeRandom.pkl", 'rb') as f:
    cartree = pickle.load(f)
dot_data = tree.export_graphviz(cartree, out_file=None, feature_names=X_train.
    ↪columns, class_names=['<50K', '>=50K'], filled=True)

# Create a graph from the Graphviz data
graph = graphviz.Source(dot_data)

# Render the graph
graph.render("decision_tree")


# Display the graph
graph
```

D:\anaconda3\lib\site-packages\sklearn\base.py:318: UserWarning: Trying to unpickle estimator DecisionTreeClassifier from version 1.2.1 when using version 1.2.2. This might lead to breaking code or invalid results. Use at your own

risk. For more info please refer to:
https://scikit-learn.org/stable/model_persistence.html#security-maintainability-limitations

```
warnings.warn(
```

[18]:




```
[19]: graph.format = 'png'  
graph.render('DecisionRandom')
```

[19]: 'DecisionRandom.png.png'

1.7 Comparing the two results

```
[24]: fig, ax = plt.subplots(1,2)  
ax[0].barh(X_train.columns,fi)  
ax[0].invert_yaxis()  
ax[0].set_title("Given Data")  
ax[1].barh(X_train.columns,fi1)  
ax[1].invert_yaxis()  
ax[1].set_title("Randomized data")
```

[24]: Text(0.5, 1.0, 'Randomized data')



```
[25]: fig.savefig('FeatureDCompare.png')
```

We can clearly observe that the feature importance for both models is quite similar

1.8 Random Forest

```
[32]: X_train.shape, X_test.shape, X_valid.shape
```

```
[32]: ((32724, 14), (8059, 14), (8059, 14))
```

```
[33]: from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(criterion='entropy', n_estimators=10)
rfc.fit(X_train, Y_train)
```

```
[33]: RandomForestClassifier(criterion='entropy', n_estimators=10)
```

```
[34]: rfc.score(X_train, Y_train)
```

```
[34]: 0.9877765554333211
```

```
[35]: rfc.score(X_valid, Y_valid)
```

```
[35]: 0.8491127931505149
```

```
[36]: rfc = RandomForestClassifier(criterion='entropy', n_estimators=30)
      rfc.fit(X_train, Y_train)
```


```
[36]: RandomForestClassifier(criterion='entropy', n_estimators=30)
```

```
[37]: rfc.score(X_train, Y_train), rfc.score(X_valid, Y_valid)
```

```
[37]: (0.9981359247035815, 0.8546966124829384)
```

```
[38]: with open("RandomForestModel.pkl", 'wb') as f:
      pickle.dump(rfc, f)
```

```
[39]: fi2 = rfc.feature_importances_
      fig, ax = plt.subplots()
      ax.barh(X_train.columns, fi2)
      ax.invert_yaxis()
```




```
[40]: fig.savefig('FeatureRF.png')
```

```
[41]: fig, ax = plt.subplots(1, 2)
      ax[0].barh(X_train.columns, fi1)
      ax[0].invert_yaxis()
      ax[0].set_title("Decision Tree")
      ax[1].barh(X_train.columns, fi2)
```



```
ax[1].invert_yaxis()
ax[1].set_title("Random Forest")
```

```
[41]: Text(0.5, 1.0, 'Random Forest')
```



```
[42]: fig.savefig('RandomDecisonCompare.png')
```

It is quite evident that the random forest sees more value in all the attributes, and is more meaningful as well from a data scientist's POV. Age can be easily understood as the factor affecting the data, while relationship being the main attribute makes little sense. Thus, with increase in attributes, decision tree becomes worse, heavier and slower.

```
[ ]: with open("RandomForestModel.pkl", 'rb') as f:
    random_forest = pickle.load(f)

for i in range(0,5):
    decision_tree = random_forest.estimators_[i]
    dot_data = tree.export_graphviz(decision_tree, out_file=None,
    ↪ feature_names=X_train.columns, class_names=['<50K', '>=50K'], filled=True)
    graph = graphviz.Source(dot_data)
    graph.format = 'png'
    graph.render('outpurandomforest/tree'+str(i))
```

1.9 Report

With this, we can finally report our findings, and complete our report. For this assignment, we used chefboost to understand the basics of C4.5 algorithm, Owing to the lack of adaptability of the algorithm, we switched to sklearn's Decision Tree Classifier, which is a CART algorithm. We have used graphviz, a tool which allowed us to plot the trees, and are saved inside the folder as .png's. For One Hot Encoding, we have used LabelEncoder of sklearn, and for splitting data, we have used traintestsplite of sklearn. We have also used Random Tree Classifier of sklearn.ensemble, and json and pickle libraries to store relevant data and model. We studied the difference between provided data and random data, and found difference to be negligible, indicating that the data is sufficiently generalised.

Then, we found the random tree classifier with 30 trees, and had a better accuracy, as well as understanding of different attributes. We plotted the Feature importance graphs of all our models, thus understanding for these models.

The Decision Tree with given data is referred to as Decision Final, random data tree as Decision Random, and Random Forest as Random Forest. We have saved the model to start working any time we require to.

1.9.1 Rules derived

For class <50k: (relationship<=0.5)^(education <=12.5) ^ (capital gain <=5095.5) ^ (education <=8.5) ^ (capital loss < = 1791.5) ^ (age <=36.5) ^ (hour per week<=49) ^ (native country <= 34.5) And so on

For class >=50k: (Capital gain >7669.5) ^ (marital status <=1) ^ (hour per week > 35.5) ^ (flwgt > 33379) ^ (age >20) ^ (education <=10.5) ^ (capital gain > 7073.5) ^ (relationship > 0.5) And so on