

Water Quality Index

Prediction and Regression Machine Learning Project

Different types of Regression Models are trained upon different forms of same dataset.

Linear Regression
Neighbors

K - Nearest

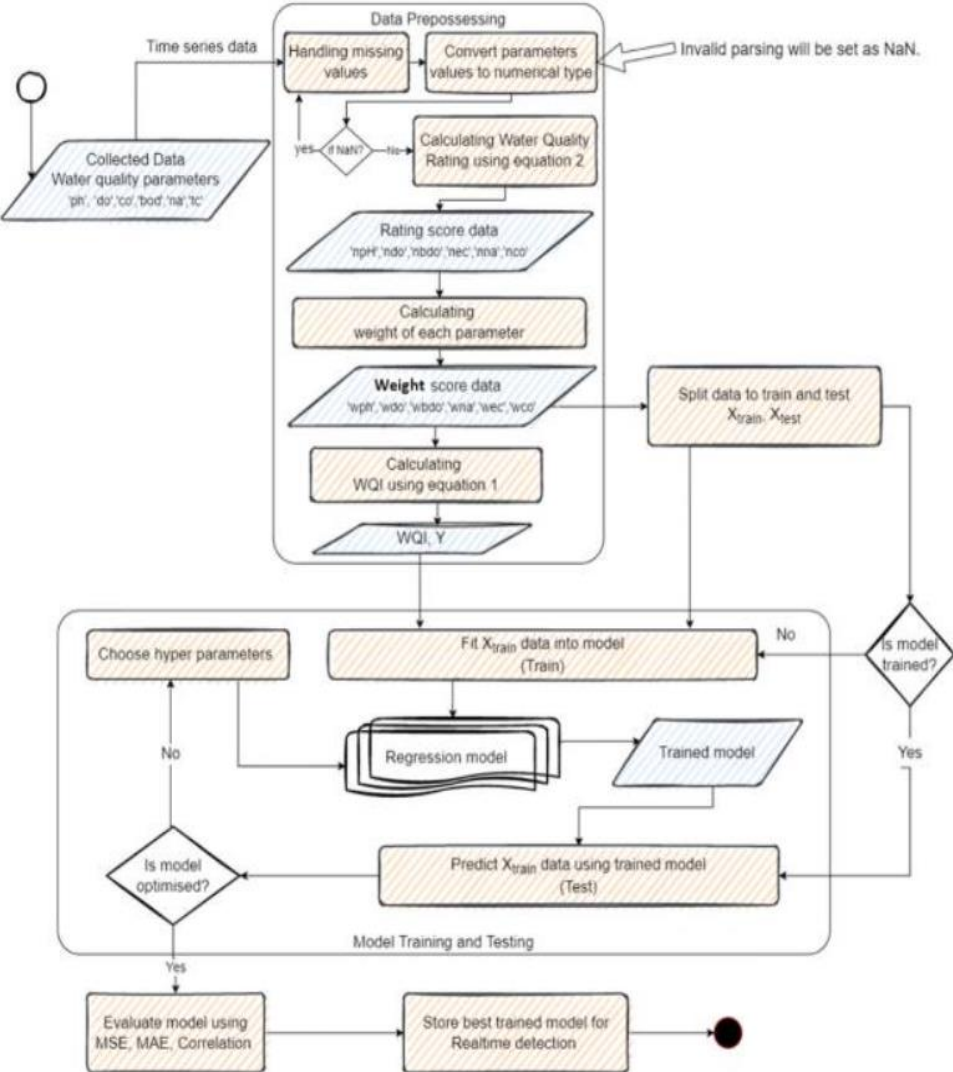
Gradient Boosting

Bayesian Ridge

Decision Tree Regression

Artificial Neural Network

Random Forest Regression



Steps involved:

- 1) Data Pre - Processing
- 2) Scaling & Normalization of values
- 3) Removing Outliers
- 4) Running the Model

```
for col in df.columns:  
    df[col].replace('?', np.NaN, inplace=True)
```


```
df['location'].fillna(df['location'].mode()[0], inplace=True)  
df['do'].fillna(df['do'].mode()[0], inplace=True)  
df['ph'].fillna(df['ph'].mode()[0], inplace=True)  
df['co'].fillna(df['co'].mode()[0], inplace=True)  
df['bod'].fillna(df['bod'].mode()[0], inplace=True)  
df['na'].fillna(df['na'].mode()[0], inplace=True)  
df['tc'].fillna(df['tc'].mode()[0], inplace=True)  
df['year'].fillna(df['year'].mode()[0], inplace=True)
```

To face the outliers present and
to prevent feature domination
and restricting the values
between 0 and 1

```
X_scaled = (X - X_min) / (X_max - X_min)
```

All missing values are filled with
NaN and then replaced with mode
of the respective columns

```
#Normalizing DataSet  
from sklearn.preprocessing import MinMaxScaler  
  
scaler = MinMaxScaler()  
  
# Fit the scaler to the data and transform the features  
X_train_scaled = scaler.fit_transform(X_train)  
  
# Transform the testing data using the fitted scaler  
X_test_scaled = scaler.transform(X_test)
```



‘Spearman’ correlation matrix

The heatmap visualization helps in understanding the strength and direction of the monotonic relationships between the variables after MinMax Scaling

```

from scipy.stats import zscore
df_num_final_norm = zscore(df_new, axis=0)

def indices_of_greater_than_3(df_norm):
    df_norm = pd.DataFrame(df_norm)
    indices_arr = []
    n_col = df_norm.shape[1]
    for index in range(n_col):
        col_index = df_norm.iloc[:, index]
        greater_than_3 = df_norm[col_index > 3]
        greater_than_3_index = greater_than_3.index
        indices_arr.extend(greater_than_3_index)
    return indices_arr

indices_arr = indices_of_greater_than_3(df_num_final_norm)
print("Number of outliers using Z-Score method-", len(indices_arr))
df_new.iloc[indices_arr, :]

```



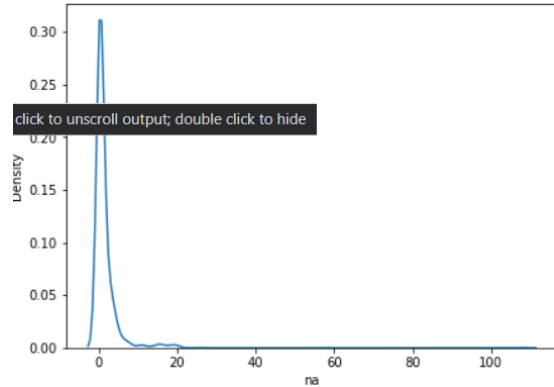



Number of outliers using Z-Score method- 114

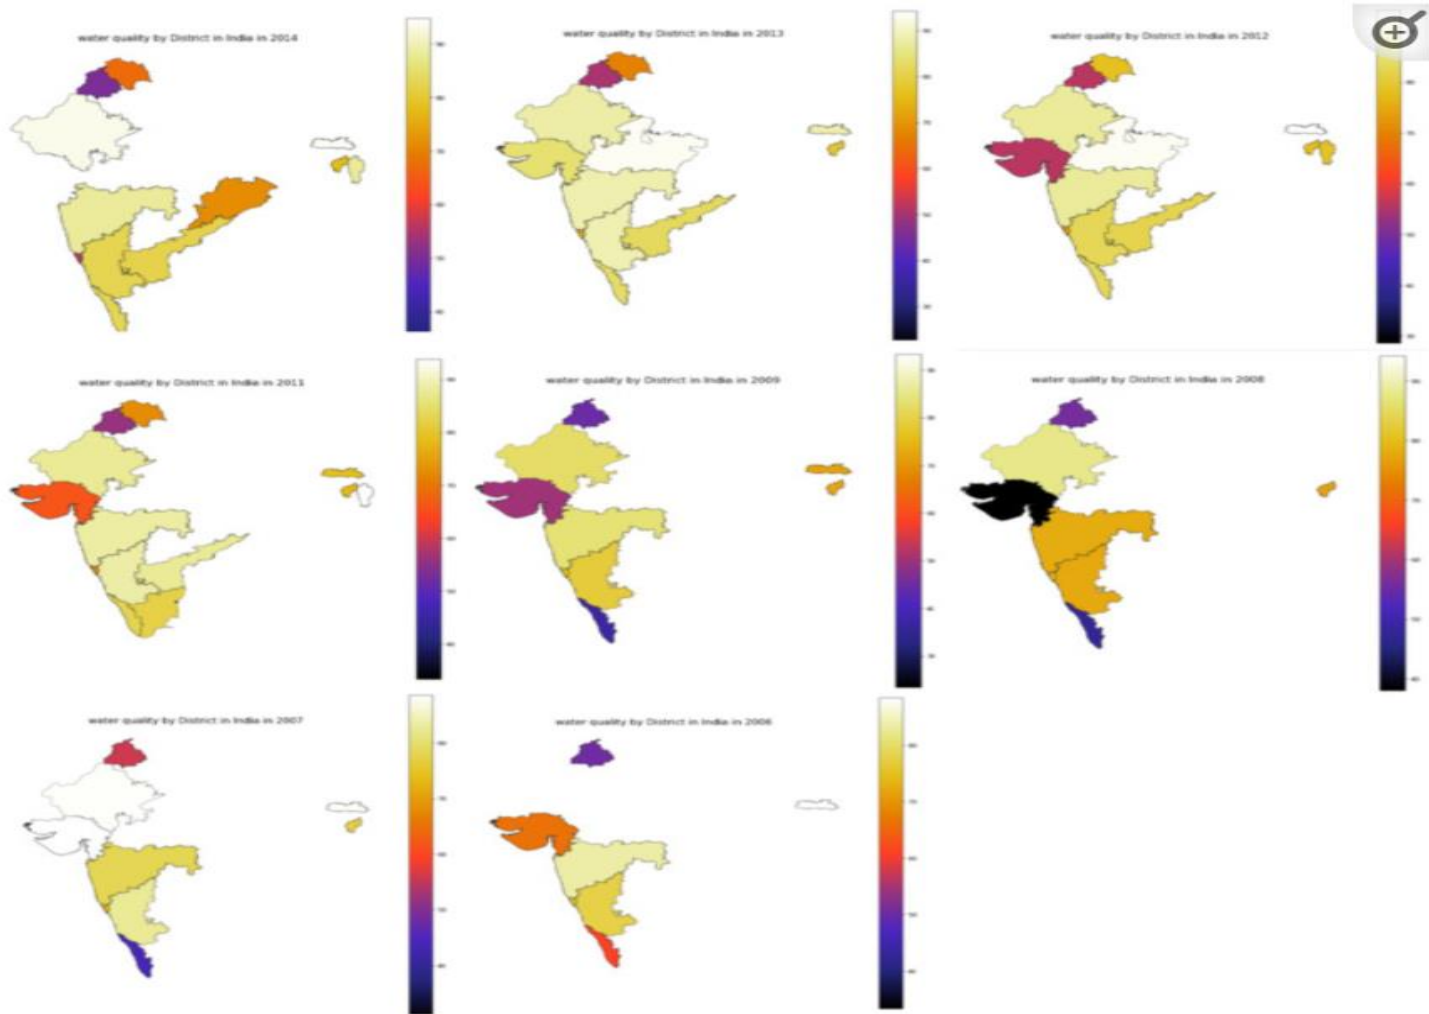
Handling outliers to remove misleading data and bias in modeling.

$$z = \frac{x - \mu}{\sigma}$$

Data points with z-scores that fall above threshold value are potential outliers

Distribution of each attribute





#Applying Gradient Boosting

```
from sklearn.ensemble import GradientBoostingRegressor
```

R2 = 0.8973223140053188

```
from sklearn.datasets import make_regression
```

```
gb = GradientBoostingRegressor(n_estimators=30, learning_rate=0.1, max_depth=6)
```

```
gb.fit(X_train_final, y_train_final)
```

```
y_predict = gb.predict(X_test_final)
```

R2 = 0.7602114513363876

```
from sklearn.tree import DecisionTreeRegressor
```

```
Treereg = DecisionTreeRegressor()
```

```
Treereg.fit(X_train_final, y_train_final)
```

```
y_pred = Treereg.predict(X_test_final)
```

```
from sklearn.ensemble import RandomForestRegressor
```

```
reg = RandomForestRegressor()
```

```
reg.fit(X_train_final, y_train_final)
```

```
y_pred = reg.predict(X_test_final)
```

R2 = 0.882445986645837

```
print(mse)
```

```
print(r2)
```

1.6635807910485109e-28

1.0

#Performing Linear Regression

```
model = LinearRegression(fit_intercept = True)
```


```
model.fit(X_train_n, y_train_n)
```

```
y_pred = model.predict(X_test_n)
```

```
mse = mean_squared_error(y_test_n, y_pred)
```

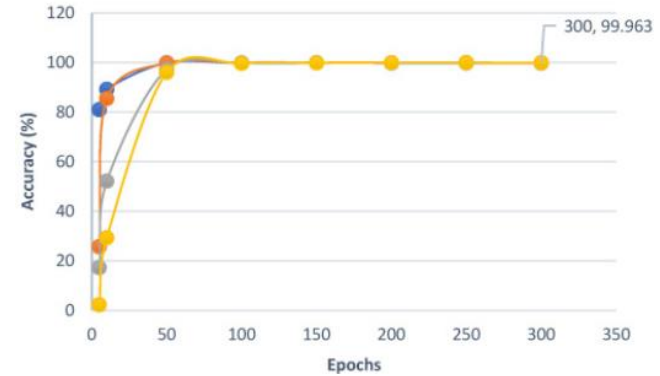
Accuracy Of Artificial Neural Network models

Model training using qi1




Qi1 = Raw Data

Model training qi2



Qi2 = Normalized Data

Model training using qi3



Qi3 = Unit Weighted Data
(using official sources)

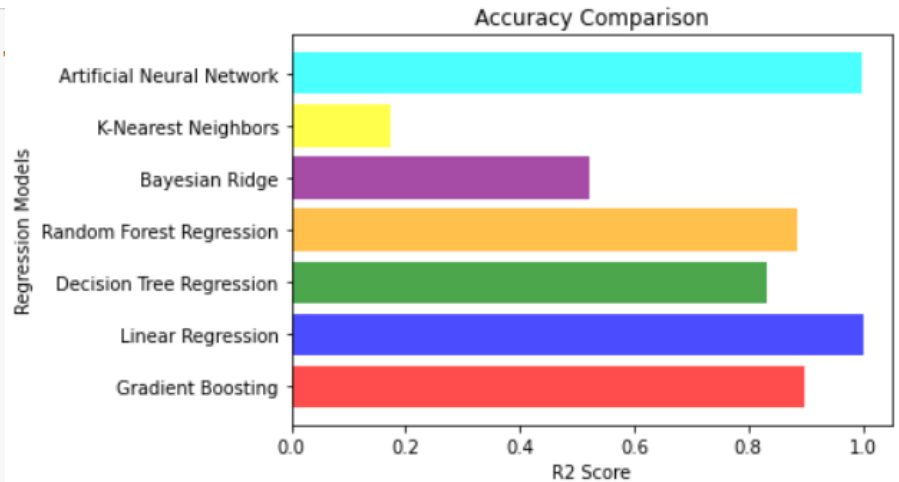
- Accuracy(Batch size: 5)
- Accuracy(Batch size: 10)
- Accuracy(Batch size: 15)
- Accuracy(Batch size: 20)

FINAL RESULTS FROM ALL MODELS USED

```
import matplotlib.pyplot as plt
labels = ['Gradient Boosting', 'Linear Regression', 'Decision Tree Regression',
          'K-Nearest Neighbors', 'Bayesian Ridge', 'Random Forest Regression', 'Artificial Neural Network']
values = [r2_gradb, r2_lrNorm, r2_tree, r2_forest, r2_br, r2_knn, r2_annNorm]
#colors = ['red', 'blue', 'green', 'orange', 'purple', 'yellow', 'cyan']
colors = [(1, 0, 0, 0.7), # Red
          (0, 0, 1, 0.7), # Blue
          (0, 0.5, 0, 0.7), # Green
          (1, 0.65, 0, 0.7), # Orange
          (0.5, 0, 0.5, 0.7), # Purple
          (1, 1, 0, 0.7), # Yellow
          (0, 1, 1, 0.7)] # Cyan

plt.barh(labels, values, color=colors)
plt.xlabel('R2 Score')
plt.ylabel('Regression Models')
plt.title('Accuracy Comparison')

plt.show()
```



Thank You!

Harshvardhan Jouhary (2020B2A31623P)

Aryaman Chauhan (2020B5A72006P)

Devansh (2020B5A72001P)